

Explanation:

“vectorizing matrix multiplication along an axis without any for loops”

$(x-\mu)$ is a `LARGE_NUMBER x 3` matrix and Σ^{-1} is `3x3` matrix. We can basically ignore the inverse part for demonstration purposes. Matrix multiplication is easy for a single row. $\Sigma^{-1} \cdot (x-\mu) = [1 \times 3] \cdot [3 \times 3]$. However, we would need to loop over every row of $(x-\mu)$ to get the full results. This is too slow. A solution I found after trying a lot of things is to add a new axis to X and do the matrix multiplication in the newly created dimension!

Basic operation:

$$\Sigma^{-1}(\underline{x} - \underline{\mu})$$

Figure 1 Σ is a `3x3` matrix $(x-\mu)$ is distance to average RGB of a cluster for a large number of independent identically distributed samples

Here is the rough process:

*also outlined and tested in `vectorized_dot_product_loop.py`:

1. Add a new axis to X of depth Σ .`rowsize`. Now X is `LARGE_NUMBER x h x 3` meaning there are 3 copies of the original matrix in the new direction added.
2. Each “row” of X is now a `3x3` matrix which can do pointwise multiplication via numpy with Σ . This is the first step of matrix multiplication if it were done manually
3. Next, you take the sum along each row to column multiplication of x and σ . This is the step that gives you the value of matrix multiplication for a row of the original X .
 - a. One note is that numpy handles axes a little funny. In the example, I add an axes on the 2nd index and sum on the 1st index. But conceptually you are creating copies of the original row in a new dimension and multiplying them by their respective column of σ and taking a sum all in the new dimension.

Take a look at the attached code if you want to play around with the idea. The pay off is that same operation over 100k rows goes from ~7s to ~0.4s when shifting from a for loop to the vectorized solution.

CONTEXT:

The scenario is implementing the expectation maximum algorithm for an unspecified number of dimensions and clusters. Basically K-Means but using a probability rather than assignment for k and an undefined number of clusters + data size. The important part for this use case is below. It is the equation for a multi-variate gaussian distribution for a sample value \underline{x} where \underline{x} is chosen from independent random samples.

$$\mathcal{N}(\underline{x} ; \underline{\mu}, \Sigma) = \frac{1}{(2\pi)^{d/2}} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (\underline{x} - \underline{\mu})^T \Sigma^{-1} (\underline{x} - \underline{\mu}) \right\}$$

For the specific project, you have a pixel (\underline{x}) with RGB values. Each cluster over a range of pixels is defined by an average ($\underline{\mu}$) of R,G, and B intensity as well as a covariance (Σ) between each value. Notably, a single pixel is 1X3. The $\underline{\mu}$ vector (average) is also 1X3 and the covariance is 3x3. This all works out nicely in the equation above and you end up with a single float value which is the probability a pixel is in the cluster defined by $\underline{\mu}$ and Σ . The problem arises when you want probability for a large number of pixels. \underline{x} becomes LARGE_NUMBER X 3 and the natural option is to loop over the large number to get your vector of probabilities. This is where the NumPy trick comes in that I think is pretty neat. After a bit of vectorizing, everything is easily moved to a single calculation for all pixels and a single cluster except for the matrix multiplication of sigma (the covariance matrix) and $\underline{x}-\underline{\mu}$ (distance of a pixel from the average RGB for a cluster). That is what is addressed via vectorizing.